User manual for R functions used in Post-Shock Persistence Testing

January 25, 2022

This document describes the functions that are used in the post-shock persistence testing.

Contents

1 scm function

`scm(X,Tstar)` is a core function in post-shock prediction methodology that is used to compute the synthetic weights $\mathbf{W}^*$ in Lin and Eck (2021). It computes the weights to minimize the Euclidean distance between the convex combination of the covariates in the donor pool and the covariates of the time series of interest. The optimization is based on nonlinear

optimization using the `solnp` in the R package `Rsolnp` (Ghalanos, Theussl, & Ghalanos, 2012).

Suppose the donor pool size is n. `scm(X,Tstar)` takes `X` and `Tstar` as inputs and outputs a vector of weights $\mathbf{W}^*$ and other convergence details related to `solnp`, where

- `X` is a list of covariates with length $n+1$, where the first element of the list should be the covariates of the time series of interest and the remaining elements of the list are the covariates of the donors. The covariates should be in the class of `matrix` with size $T_i \times p$ for $i = 1, \ldots, n+1$.

- `Tstar` is a vector of shock-effect time points when the shock occurs at $T_i^* + 1$ for $i = 1, \ldots, n+1$. The size of `Tstar` should be $n+1$.

- The default for `scale` is `scale = FALSE`. If `scale = TRUE`, the function collects $\mathcal{Q} = \{\mathbf{x}_{i,T_i^*+1} : i = 1, \ldots, n+1\}$ and scale and center within $\mathcal{Q}$, i.e., center by the sample mean of $\{x_{i,j,T_i^*+1} : i = 1, \ldots, n+1\}$, where j index stands for the jth variable.

Examples:

```
p <- 2; n <- 10;
T <- round(rgamma(n = n + 1, shape = 15, scale = 10)) # Time Length
Tstar <- c() # Shock Time Points
for (t in T) {
  Tstar <- c(Tstar, sample((p + 3 + 1):(t - 1), size = 1))
}
phi <- round(runif(n + 1, 0, 1), 3) # autoregressive parameters
# construction of design matrix
X <- c()
for (i in 1:(n + 1)) {
  Ti <- T[i]
  Tstari <- Tstar[i]
  X[[i]] <- matrix(rgamma(n = p * (Ti + 1), shape = 1, scale = 1),
                   ncol = p, byrow = T)
}
scm(X, Tstar)
```

2 `sel` function

`sel(yhat,y)` function is a function that computes the squared error loss between $\hat{y}$ and y, where

- `yhat` is either a vector or a scalar

- `y` is either a vector or a scalar

`yhat` and `y` must have equal length. The output of `sel(yhat,y)` is the squared error loss. It may be a vector or scalar.

Examples:

```
sel(rnorm(10), rep(0,10))
```

3   `taSPA.mhfc` function

`taSPA.mhfc(ell,d,B,bw)` is a core function that implements the hypothesis test of Quaedvlieg (2021) that compares two forecasts in terms of average superior predictive ability (aSPA) for multiple horizons. See the theoretical and algorithmic details of the hypotheses in Quaedvlieg (2021). Let H be the maximum horizon number for the forecasts and T be the sample size. The test statistic is

$$t_{aSPA,ij} = \frac{\sqrt{T}\bar{d}_{ij}}{\hat{\zeta}_{ij}},$$

where $d_{ij,t}$ ($H \times 1$) is the loss differential between forecasts i and j at time t, $\bar{d}_{ij}$ is the average across $t = 1, \ldots, T$, $\bar{d}_{ij} = w'\bar{d}_{ij}$, $\zeta_{ij} = \sqrt{w'\Omega_{ij}w}$ is estimated by HAC estimator, Ω_{ij} is the asymptotic variance of $\bar{d}_{ij}$. $w = \mathbf{1}_H/H$ is taken as default. The computation of HAC estimator requires specification of a bandwidth parameter `bw`. Provided with a loss differential matrix `d`, this test uses moving-block bootstrap to approximate the distribution of the test statistic. Note that

- `ell` is the block length of the moving block bootstrap. Theory of Quaedvlieg (2021) requires $ell = o(T^{1/2})$.

- `d` is the loss differential matrix of size $T \times H$. It must be of `matrix` class. The element in tth row and hth column should be the loss differential between two h-step forecasts for tth observation.

- `B` is the bootstrap sample size.

- `bw` is the bandwidth of computation for HAC estimator. HAC estimation is based on `HAC` from Schlittgen (2021) using quadratic spectral kernel.

The output of `taSPA.mhfc(ell,d,B,bw)` is the *p*-value of the test, which is a scalar. Importantly, `taSPA.mhfc(ell,d,B,bw)` tests a *one-sided* hypothesis. Thus it is necessary to check whether the way about how the loss differential `d` is computed matches the hypothesis that is tested. See more details of the hypothesis that is tested in Quaedvlieg (2021).

Examples:

```
y <- arima.sim(list(2,0,0), n = 100)
```

```
T <- 50; H <- 10; K <- 40

m1.L.i <- matrix(NA, nrow = T, ncol = H)
m2.L.i <- matrix(NA, nrow = T, ncol = H)
for (h in 1:H) {
  for (t in 1:T) {
    y.t <- y[(t + H - h + 1):(t + K + H - h)]

    m1 <- arima(y.t, order = c(1, 0, 0))
    m2 <- arima(y.t, order = c(2, 0, 0))

    y.t.hat.1 <- predict(m1, h)$pred[h]
    y.t.hat.2 <- predict(m1, h)$pred[h]

    m1.L.i[t, h] <- sel(y = y[t + K + H], yhat = y.t.hat.1)
    m2.L.i[t, h] <- sel(y = y[t + K + H], yhat = y.t.hat.2)
  }
}
d <- m2.L.i - m1.L.i

taSPA.mhfc(ell = 2, d = d, B = 200, bw = 4)
```

4 `vote` function

`vote(ps,sig)` is a function that yields the significance voting from a vector of p-values given a significance level. To be specific, if over 50% of the tests, each of which corresponds to a p-value in the donor pool, are significant according to the given significance level, the output of `vote(ps,sig)` would be 1. It is 0 otherwise. Note that

- `ps` is a vector of p-values with values in $[0, 1]$.

- `sig` is the significance level. It is 0.05 in default. It should be a scalar.

  Examples:
```
vote(runif(10), .05)
```

5 `ols.est.alphahat` function

`ols.est.alphahat(Tstar,Y,X)` is the core function in Lin and Eck (2021) that is used to compute the weighted adjusted shock effect estimate $\hat{\alpha}_{\text{wadj}}$ that aggregates information in the donor pool. Suppose the donor pool size is n. Given the shock time points, it performs n OLS regressions across the donor pool to estimate the shock effect of each donor.

The regression is based on the **AR**(1) model with covariates as specified in Section 2.1 of Lin and Eck (2021). The function internally uses `scm` function to compute the weights $\mathbf{W}^*$ and weight the donors' shock effects to compute $\hat{\alpha}_{\mathrm{wadj}}$. It also computes the $\hat{\alpha}_{\mathrm{IVW}}$ estimate and $\hat{\alpha}_{\mathrm{adj}}$ estimate. See more details of these two estimates in Lin and Eck (2021). Note that

- `X` is a list of covariates with length $n + 1$, where the first element of the list should be the covariates of the time series of interest and the remaining elements of the list are the covariates of the donors. The covariates should be in the class of `matrix` with size $T_i \times p$ for $i = 1, \ldots, n + 1$.

- `Tstar` is a vector of shock-effect time points when the shock occurs at $T_i^* + 2$ for $i = 1, \ldots, n + 1$. The size of `Tstar` should be $n + 1$.

- `Y` is a list of responses with length $n + 1$, where the first element of the list should be the response of the time series of interest and the remaining elements of the list are the responses of the donors. It should be a vector. Note that to compute the lag, the effective sample size used is $T_i - 1$, where T_i is the time series length of ith time series.

This function outputs

- `alphahat` is a vector of n length and contains the OLS shock effect estimates for the donors.

- `est` is a vector of three shock effect estimates, $\hat{\alpha}_{\mathrm{adj}}$, $\hat{\alpha}_{\mathrm{wadj}}$, and $\hat{\alpha}_{\mathrm{IVW}}$.

- `Tstar` is the same as the one in the input.

- `X` is the same as the one in the input.

- `Y` is the same as the one in the input.

- `lmod` is a list of n `lm` objects that were fitted by OLS for each donor. The ordering is the same as how the donor is ordered in `X`, `Y`, and `Tstar`.

- `res` is a list of n residual objects corresponding to each element of `lmod`.

- `Wstar` is the synthetic weights, as defined in the output of `scm`.

- `se` is a vector of length n and each element is the OLS standard error for the OLS shock effect estimate for each donor.

Examples:

```
n <- 10; p <- 2
Ts <- ceiling(rgamma(n + 1, scale = 10, shape = 10)) # Time Length
Tstar <- c() # Shock Time Points
for (t in Ts) {
  Tstar <- c(Tstar, sample((p + 3 + 1):(t - 1), size = 1))
}
phi <- round(runif(n + 1, 0, 1), 3) # autoregressive parameters
```

```
X <- c()
alpha <- c()
# construction of design matrix and shock effects
gamma <- c()
for (i in 1:(n + 1)) {
  Ti <- Ts[i]
  Tstari <- Tstar[i]
  X[[i]] <- matrix(rgamma(n = p * (Ti + 1),
                          shape = 1, scale = 1), ncol = p, byrow = TRUE)
  gamma[[i]] <- matrix(rnorm(p, mean = 1, sd = 1), nrow = 1)
  epsilontildei <- rnorm(n = 1, sd = 1)
  alpha <- c(alpha, 1 + gamma[[i]] %*% X[[i]][Tstari + 1, ] +
               epsilontildei)
}

# generation of yit
Y <- c()
for (i in 1:(n + 1)) {
  yi0 <- rnorm(1)
  Tstari <- Tstar[i]
  alphai <- alpha[i]
  phii <- phi[i]
  xi <- X[[i]]
  yi <- yi0
  thetai <- matrix(rnorm(p), nrow = 1)
  etai <- rnorm(1)
  for (t in 2:(Ts[i] + 1)) {
    epsilonit <- rnorm(n = 1, sd = 1)
    yi <- c(yi, etai + alphai * ifelse(t >= Tstari + 2,
                                       yes = 1, no = 0) +
            phii * yi[t - 1] + thetai %*% xi[t, ] + epsilonit)
  }
  Y[[i]] <- yi
}
ols.est.alphahat(Tstar,Y,X)
```

6 `ps.indic.W.permanent`

`ps.indic.W.permanent(Tstar,Y,X,K,H,Ts,ell,B,bw,sig.levl)` is a core function that produces the synthetic weights from `scm`, a sequence of p-values and corresponding rejection/acceptance decisions for each donor in the donor pool based on a give significance level and the test function `taSPA.mhfc`. The assumed model for this function

is

$$y_{i,t} = \eta_i + \phi_i y_{i,t-1} + x_{i,t}\theta + \alpha_i I(t > T_i^* + 1) + \varepsilon_{i,t},$$
$$\alpha_i = \mu_\alpha + x_{i,T_i^*+1} + \varepsilon_{\alpha,i},$$

where $T_i^* + 1$ is the time point when the shock occurs for $i = 2, \ldots, n+1$, $x_{i,t}$ is the covariate, and η_i is the intercept. Note that only θ, η_i, and α_i are estimable. The above model is called permanent model due to the presence of a persistent shock term. For each donor, this function computes the adjusted forecast and unadjusted forecast, construct the corresponding loss differential matrix, and perform the forecast comparison test. Suppose the donor pool size is n. *It is important to note that the ordering of the donors should be with respect to the timing when those time series occur, from the earliest to the latest.*

The input for this function is

- `Tstar` is a vector of shock-effect time points when the shock occurs at $T_i^* + 2$ for $i = 1, \ldots, n+1$. The size of `Tstar` should be $n+1$.

- `Y` is a list of responses with length $n+1$, where the first element of the list should be the response of the time series of interest and the remaining elements of the list are the responses of the donors. It should be a vector. Note that to compute the lag, the effective sample size used is $T_i - 1$, where T_i is the time series length of ith time series.

- `X` is a list of covariates with length $n+1$, where the first element of the list should be the covariates of the time series of interest and the remaining elements of the list are the covariates of the donors. The covariates should be in the class of `matrix` with size $T_i \times p$ for $i = 1, \ldots, n+1$. Note that the sample size of each element in `X` should be $T_i + K_i + H$, where K_i is the training sample size and T_i is the length of the donor time series that are wished to be considered as the data used in forecast comparison. H is the maximum horizon.

- `K` is a vector of training sample sizes that are used to compute h-step forecast for each data point. `K` has a length of $n+1$. If `retro = FALSE`, `K[1]` can be `NA`.

- `H` is the maximum h that an h-step forecast that is wished to be considered in the multi-horizon forecast comparison testing procedure in Quaedvlieg (2021).

- `Ts` is a vector of length $n+1$, and it specifies the length of the time series (ordered from the time series of interest to donors). Note that for each element of `Ts`, it should be at least greater than or equal to $K_i + H + 2$.

- `ell` is the block length of the moving block bootstrap. See details in `taSPA.mhfc`.

- `B` is the bootstrap sample size. See details in `taSPA.mhfc`.

- `bw` is the bandwidth of computation for HAC estimator. See details in `taSPA.mhfc`.

- `sig.levl` is the significance level for the test. It is 0.05 in default.

- `retro` is a logical value indicating `TRUE` or `FALSE`. If `TRUE`, the p-value for the time series of interest will be computed. If `FALSE`, only the p-values of donors will be computed. This argument is intended for retrospective data analysis where the p-value of the time series of interest is needed.

- `scale` option can allow users to compute weights based on scaled covariates or not. See details in `scm`.

- `sc.x` is a list of one vector and a list, where the ordering matters. The first vector should be a vector of indices referring to which time series wish to consider additional covariates that are present in the models but not considered in the weight computation. The second list should be a list of matrices with length equal to the length of the first vector. The matrices are the matrices of covariates, which are present in the models but not considered in the weight computation.

The output of this function is

- `ps` is a sequence of p-values from forecast comparison test of Quaedvlieg (2021) for the donors. It is of length n.

- `W` is a vector of synthetic weights and is of length n.

- `Is` is a sequence of rejection/acceptance decisions from forecast comparison for the donors, where 1 stands for rejection and 0 for acceptance. It is of length n.

Examples:

```
n <- 10; p <- 2
K <- ceiling(rgamma(n + 1, scale = 5, shape = 10)) # training sample size
Ts <- ceiling(rgamma(n + 1, scale = 10, shape = 10)) # Time Length
Tstar <- c()
for (i in 1:(n + 1)) {
  Tstar <- c(Tstar,  max(Ts[i] + 1, ceiling(0.5 * (Ts[i] + K[i] + H))))
}
phi <- round(runif(n + 1, 0, 1), 3) # autoregressive parameters

X <- c()
alpha <- c()
# construction of design matrix and shock effects
gamma <- c()
for (i in 1:(n + 1)) {
  Ti <- Ts[i]
  Tstari <- Tstar[i]
  Ki <- K[i]
  X[[i]] <- matrix(rgamma(n = p * (Ti + Ki + H + 1),
                          shape = 1, scale = 1), ncol = p, byrow = TRUE)
  gamma[[i]] <- matrix(rnorm(p, mean = 1, sd = 1), nrow = 1)
  epsilontildei <- rnorm(n = 1, sd = 1)
```

```
   alpha <- c(alpha, 1 + gamma[[i]] %*% X[[i]][Tstari + 1, ] +
                 epsilontildei)
}

# generation of yit
Y <- c()
for (i in 1:(n + 1)) {
  yi0 <- rnorm(1)
  Tstari <- Tstar[i]
  alphai <- alpha[i]
  phii <- phi[i]
  xi <- X[[i]]
  yi <- yi0
  thetai <- matrix(rnorm(p), nrow = 1)
  etai <- rnorm(1)
  for (t in 2:(K[i] + Ts[i] + H + 1)) {
    epsilonit <- rnorm(n = 1, sd = 1)
    yi <- c(yi, etai + alphai * ifelse(t >= Tstari + 2,
                                       yes = 1, no = 0) +
            phii * yi[t - 1] + thetai %*% xi[t, ] + epsilonit)
  }
  Y[[i]] <- yi
}

est <- ps.indic.W.permanent(Tstar = Tstar, Y = Y, X = X, K = K,
                            H = 10, Ts = Ts, ell = 4, B = 200, bw = 4)
```

7 `ps.indic.W.dynamic`

`ps.indic.W.dynamic(Tstar,Y,X,K,H,Ts,q1,q2,ell,B,bw,sig.levl)` performs
the similar functionality as `ps.indic.W.permanent` but with a different model as

$$
y_{i,t} = \left(\eta_i + \sum_{j=1}^{q_1} \phi_{i,j} y_{i,t-j} + \sum_{j=0}^{q_2-1} x_{i,t-j} \theta_{i,j+1} \right) (1 - D_{i,t}) + f(\mathcal{F}_{i,t}, \alpha_i) D_{i,t} + \varepsilon_{i,t}
$$

$$
f(\mathcal{F}_{i,t}, \alpha_i) = \alpha_i + \sum_{j=1}^{q_1} \tilde{\phi}_{i,j} y_{i,t-j} + \sum_{j=0}^{q_2-1} x_{i,t-j} \tilde{\theta}_{i,j+1}
$$

$$
\alpha_i = \mu_\alpha + x_{i,T_i^*+1} + \varepsilon_{\alpha,i}
$$

where $D_{i,t} = I(t > T_i^* + 1)$. Note that only $\eta_i, \alpha_i - \eta_i, \phi_{i,j}, \theta_{i,j+1}, \tilde{\phi}_{i,j} - \phi_{i,j}, \tilde{\theta}_{i,j+1} - \theta_{i,j+1}$ are estimable. The role of the shock in this model is that the dynamics of $y_{i,t}$ change after the shock. Note that to prepare the lagged response and covariates, the effective sample size is

9

$T_i - \max\{q_1, q_2 - 1\}$. Note that for each element of `Ts`, it should be at least greater than or equal to $K_i + H + \max\{q_1, q_2 - 1\} + 1$.

Additional functionality of `ps.indic.W.dynamic` compared to `ps.indic.W.permanent` is the specification of `nolag.i.x`, which allows users to add extra covariates that are not wished to be lagged in the model of $y_{i,t}$. `nolag.i.x` should be a list object of length two with the first element being a vector of indices referencing to the donor. The second element should be a list object with elements being the added covariates (of matrix or vector type) that are not wished to be lagged.

Examples are similar to those for `ps.indic.W.permanent`.

References

Ghalanos, A., Theussl, S., & Ghalanos, M. A. (2012). Package âĂŸrsolnpâĂŹ.

Lin, J., & Eck, D. J. (2021). Minimizing post-shock forecasting error through aggregation of outside information. *International Journal of Forecasting*.

Quaedvlieg, R. (2021). Multi-horizon forecast comparison. *Journal of Business & Economic Statistics*, *39*(1), 40–53.

Schlittgen, R. (2021). R package 'tsapp'.